

# 1. Undefined Behavior as a Domain Restriction

## 1.1. TL;DR

**Undefined behavior is not “wrong output.”**

It means the program has stepped *outside the compiler’s contract*.

Unreachable UB is allowed; reachable UB voids all guarantees.

This document tries to capture the Rektoff Cohort 3, office hours conversation around Rust’s undefined behavior.

## 1.2. Model

Let:

- $P :=$  the set of all **syntactically valid** Rust programs
- $S \subset P :=$  programs for which **no possible execution** can trigger UB
- $U := P \setminus S$
- $C :=$  the set of machine-code executables

Model the compiler as a **partial function**:

$\text{compile} : S \rightarrow C$

It is defined only on  $S$ . Programs in  $U$  lie outside the compiler’s domain of definition.

*The compiler maps valid Rust programs to well-defined executables.*

## 1.3. Reachability Is the Boundary

Membership in  $S$  is a property of **executions** ( $e$ ), not syntax.

Formally:

$$p \in S \quad \text{if} \quad \forall e \in \text{Exec}(p), e \text{ does not trigger UB}$$

Unreachable undefined behavior is permitted; reachable undefined behavior is not.

## 1.4. Why This Matters

### 1.4.1. Unreachable UB: allowed (in $S$ )

```
if false {
    unsafe {
        *(0 as *mut u8) = 1;
    }
}
```

- Syntactically valid (in  $P$ )
- The branch is provably unreachable
- No execution can trigger UB
- Therefore the program is in  $S$ , and compile is well-defined

### 1.4.2. Reachable UB: forbidden (in $U$ )


```
if cond { // get rekt!
    unsafe {
        *(0 as *mut u8) = 1;
    }
}
```

- There exists an execution where `cond == true`
- UB is reachable
- Therefore the program is in  $U$

- The compiler's semantic guarantees no longer apply

This holds even if cond is “always false in practice”.

## 1.5. Diagram (Mental Model)



Only programs inside  $S$  are valid inputs to the compiler-as-a-function.

## 1.6. Consequences

$\text{UB} \neq \text{incorrect output}$

$\text{UB} = \text{execution outside the compiler's contract}$

If a program admits an execution that triggers UB:

- The program lies outside the compiler's domain of definition
- The compiler always assumes its input satisfied Rust's rules
- No semantic guarantees apply to the resulting behavior

## 1.7. Contrast: Rust vs C/C++

Key difference: where the boundary is enforced

### Rust

- UB is defined in terms of possible executions
- Unreachable UB is explicitly allowed
- Safety rules are designed to preserve soundness under aggressive optimization (the whole pipeline of AST -> MIR -> ... LLVM) operates on trusted inputs. GIGO caveats.

### C / C++

- UB is often tied more loosely to syntax and local reasoning
- “This code never runs” arguments are more fragile. (Oh yea? prove it!)
- Compilers may generate UB in ways that surprise even experienced developers

Put differently:

Rust draws a hard semantic boundary around what may execute. C and C++ often leave that boundary implicit, and leaky.

## 1.8. Conclusion

A Rust program is valid only if every possible execution stays within the compiler's assumptions.

## 1.9. The obvious lingering question:

- How do you prove a rust program has no UB?